Final Project:
Form a group of 3 or 4 students. Choose a dataset from the UC Irvine Machine Learning Repository (https://archive.ics.uci.edu/ml/index.php) with at least 5000 instances and 20 attributes for classification or regression. Compare how the different approaches seen in class perform on this dataset to predict accurately the classes or the values of the unlabeled data. You should determine what are the best hyper-parameters for each approach you are using. You could use any Python libraries.
You must write a paper (8 to 10 pages) about your project. You must include a presentation of the research questions and the chosen methods to tackle them, Relevant Literature Review, a presentation of the results and discussion and a conclusion/future work.

This course assumes the student will allocate a significant amount of independent study and time spent on coding programs and other AI simulation tools. You are strongly encouraged to ensure sufficient time needed to succeed in this course.